

practical2

May 4, 2024

```
[1]: # Data Wrangling II
# Create an "Academic performance" dataset of students and perform the
# following operations using
# Python.
# 1. Scan all variables for missing values and inconsistencies. If there are
# missing values and/or
# inconsistencies, use any of the suitable techniques to deal with them.
# 2. Scan all numeric variables for outliers. If there are outliers, use any
# of the suitable techniques
# to deal with them.
# 3. Apply data transformations on at least one of the variables. The purpose
# of this
# transformation should be one of the following reasons: to change the
# scale for better
# understanding of the variable, to convert a non-linear relation into a
# linear one, or to decrease
# the skewness and convert the distribution into a normal distribution.
```

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from sklearn.preprocessing import StandardScaler
```

```
[3]: data = pd.read_csv('../academic performance.csv')
data.head()
```

```
[3]:  gender NationalITy PlaceofBirth    StageID GradeID SectionID Topic \
0      M             KW      KuwaIT  lowerlevel    G-04         A    IT
1      M             KW      KuwaIT  lowerlevel    G-04         A    IT
2      M             KW      KuwaIT  lowerlevel    G-04         A    IT
3      M             KW      KuwaIT  lowerlevel    G-04         A    IT
4      M             KW      KuwaIT  lowerlevel    G-04         A    IT

Semester Relation  raisedhands  VisITedResources  AnnouncementsView \
```

0	F	Father	15.0	16	2
1	F	Father	20.0	20	3
2	F	Father	10.0	7	0
3	F	Father	30.0	25	5
4	F	Father	40.0	50	12

	Discussion	ParentAnsweringSurvey	ParentschoolSatisfaction	\
0	20	Yes	Good	
1	25	Yes	Good	
2	30	No	Bad	
3	35	No	Bad	
4	50	No	Bad	

	StudentAbsenceDays	Class
0	Under-7	M
1	Under-7	M
2	Above-7	L
3	Above-7	L
4	Above-7	M

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 480 entries, 0 to 479
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 480 non-null    object
1   NationalITy            480 non-null    object
2   PlaceofBirth           480 non-null    object
3   StageID                480 non-null    object
4   GradeID                480 non-null    object
5   SectionID              480 non-null    object
6   Topic                  480 non-null    object
7   Semester               480 non-null    object
8   Relation               480 non-null    object
9   raisedhands            455 non-null    float64
10  VisITedResources       480 non-null    int64
11  AnnouncementsView      480 non-null    int64
12  Discussion              480 non-null    int64
13  ParentAnsweringSurvey  480 non-null    object
14  ParentschoolSatisfaction 480 non-null    object
15  StudentAbsenceDays     480 non-null    object
16  Class                  480 non-null    object
dtypes: float64(1), int64(3), object(13)
memory usage: 63.9+ KB
```

```
[5]: # 1. Scan all variables for missing values and inconsistencies. If there are
      ↪missing values and/or
      # inconsistencies, use any of the suitable techniques to deal with them.
      data.isnull().sum() #check for missing values
```

```
[5]: gender                0
      NationalITy          0
      PlaceofBirth         0
      StageID              0
      GradeID              0
      SectionID            0
      Topic                0
      Semester             0
      Relation             0
      raisedhands          25
      VisITedResources     0
      AnnouncementsView    0
      Discussion           0
      ParentAnsweringSurvey 0
      ParentschoolSatisfaction 0
      StudentAbsenceDays   0
      Class                0
      dtype: int64
```

```
[6]: # we can remove null values by dropping rows or fill null values as per features
      # data1 = data.interpolate()
```

```
[7]: # data1.isnull().sum()
```

```
[8]: mean_val = data['raisedhands'].mean()
      data['raisehands']=data['raisehands'].fillna(mean_val)
      data2.isnull().sum()
```

```
[8]: gender                0
      NationalITy          0
      PlaceofBirth         0
      StageID              0
      GradeID              0
      SectionID            0
      Topic                0
      Semester             0
      Relation             0
      raisedhands          0
      VisITedResources     0
      AnnouncementsView    0
      Discussion           0
      ParentAnsweringSurvey 0
```


```
ParentschoolSatisfaction    0
StudentAbsenceDays          0
Class                       0
dtype: int64
```

```
[9]: # Or we can drop null value rows
      # data.dropna(axis=0, inplace=True)
```

```
[10]: # 2. Scan all numeric variables for outliers. If there are outliers, use any
      # of the suitable techniques
      # to deal with them.
fig,axis = plt.subplots(1, 2)
max_val = data2.VisITedResources.quantile(0.95)
min_val = data2.VisITedResources.quantile(0.05)
print("Before Shape", data.shape)
df = data2[(data2['VisITedResources'] > min_val) & (data2['VisITedResources'] <
↳max_val)]
print("After Shape", df.shape)
sns.boxplot(data2['VisITedResources'], orient='v', ax=axis[0])
axis[0].title.set_text('Before')
sns.boxplot(df['VisITedResources'], orient='v', ax=axis[1])
axis[1].title.set_text('After')
plt.show()
```

Before Shape (480, 17)

After Shape (427, 17)




```
[11]: # 3. Apply data transformations on at least one of the variables. The purpose
      # of this
      # transformation should be one of the following reasons: to change the
      # scale for better
      # understanding of the variable, to convert a non-linear relation into a
      # linear one, or to decrease
      # the skewness and convert the distribution into a normal distribution.
```

```
[12]: scaler = StandardScaler()
x = df[['raisedhands', 'VisITedResources', 'AnnouncementsView', 'Discussion']]
scaledf = scaler.fit_transform(x)
print(scaledf)
```


```
[[-1.15922617 -1.29844429 -1.42428794 -0.90164582]
 [-0.99010159 -1.16788899 -1.38643105 -0.71803137]
 [-1.32835075 -1.59219371 -1.50000172 -0.53441693]
 ...
 [ 0.19377045  0.59460753 -0.55357945 -0.57113982]
 [-0.65185244 -1.26580546 -0.97000525  0.45710106]
 [-0.48272786 -1.36372194 -0.62929323  0.6407155 ]]
```

```
[13]: df.Discussion.hist()  
plt.show()
```



```
[14]: import scipy.stats as stats
```

```
[15]: sns.distplot(df['Discussion'], bins = 40)  
plt.show()
```



```
[16]: # Checking the skewness  
df['Discussion'].skew()
```

```
[16]: 0.33203952447202845
```

```
[17]: # If you want to reduce skewness there are 4 methods, one of them is log  
log = np.log(df['Discussion'])  
print(log.skew())
```

```
-1.1910468123704019
```

```
[18]: sns.distplot(log, bins = 40)
```

```
[18]: <Axes: xlabel='Discussion', ylabel='Density'>
```

