

untitled14

May 3, 2024

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
[2]: data=pd.read_csv("C:\\Users\\nayan\\Downloads\\HousingData.csv")
data
```

```
[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
..
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	NaN	36.2
..
501	21.0	391.99	NaN	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

[506 rows x 14 columns]

```
[3]: data.dropna(how="any", inplace=True)
```

```
[4]: data.isnull().sum()
```

```
[4]: CRIM      0
      ZN       0
      INDUS   0
      CHAS    0
      NOX     0
      RM      0
      AGE     0
      DIS     0
      RAD     0
      TAX     0
      PTRATIO 0
      B       0
      LSTAT   0
      MEDV    0
      dtype: int64
```

```
[5]: data.columns
```

```
[5]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
          'PTRATIO', 'B', 'LSTAT', 'MEDV'],
          dtype='object')
```

```
[6]: x=data[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
          'PTRATIO', 'B', 'LSTAT']]
```

```
[7]: x
```

```
[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3	222	
..	
499	0.17783	0.0	9.69	0.0	0.585	5.569	73.5	2.3999	6	391	
500	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6	391	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
	PTRATIO	B	LSTAT								
0	15.3	396.90	4.98								
1	17.8	396.90	9.14								

```

2      17.8  392.83  4.03
3      18.7  394.63  2.94
5      18.7  394.12  5.21
..      ...      ...      ...
499    19.2  395.77  15.10
500    19.2  396.90  14.33
502    21.0  396.90  9.08
503    21.0  396.90  5.64
504    21.0  393.45  6.48

```

[394 rows x 13 columns]

```
[8]: y=data['MEDV']
```

```
[9]: y
```

```

[9]: 0      24.0
     1      21.6
     2      34.7
     3      33.4
     5      28.7
     ...
     499    17.5
     500    16.8
     502    20.6
     503    23.9
     504    22.0
     Name: MEDV, Length: 394, dtype: float64

```

```

[10]: ## Splitting train and test data using scalar
      x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.25,
      ↪random_state=42)

```

```

[11]: model = LinearRegression()
      model.fit(x_train, y_train)

```

```
[11]: LinearRegression()
```

```

[12]: #predict
      y_pred=model.predict(x_test)

```

```
[13]: model.score(x_train, y_train)
```

```
[13]: 0.7894375551691273
```

```
[14]: model.score(x_test, y_test)
```

```
[14]: 0.6830391070219286
```

```
[15]: np.sqrt(mean_squared_error(y_test, y_pred))
```

```
[15]: 5.4572694398436115
```

```
[ ]:
```